


# GREBE: Unveiling Exploitation Potential for Linux Kernel Bugs

Zhenpeng Lin, Yueqi Chen, Yuhang Wu, Dongliang Mu,  
Chensheng Yu, Xinyu Xing, Kang li

2022 IEEE Symposium on Security and Privacy (SP)

# Background

```
1 static void tun_attach(struct tun_struct *tun, ...)
2 {
3     if (tun->flags & IFF_NAPI) {
4         // initialize a timer
5         hrtimer_init(&napi->timer, CLOCK_MONOTONIC,
6                     HRTIMER_MODE_REL_PINNED);
7         // link current napi to the device's napi list
8         list_add(&napi->dev_list, &dev->napi_list);
9     }
10 }
11
12 static void tun_detach(struct tun_file *tfile, ...)
13 {
14     struct tun_struct *tun = rtnl_dereference(tfile->tun);
15     if (tun->flags & IFF_NAPI) {
16         // GPF happens if timer is uninitialized
17         hrtimer_cancel(&tfile->napi->timer);
18         // remove the current napi from the list
19         netif_napi_del(&tfile->napi);
20     }
21     destroy(tfile); // free napi
22 }
23
24 void free_netdev(struct net_device *dev) {
25     list_for_each_entry_safe(p, n,
26                             &dev->napi_list, dev_list)
27         netif_napi_del(p); // use-after-free
28 }
```

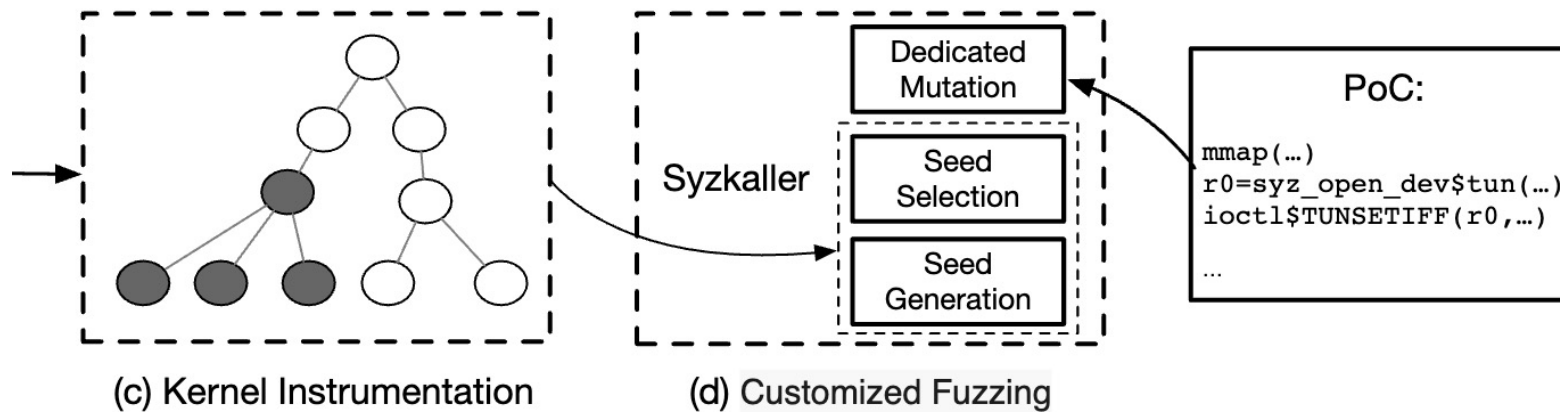
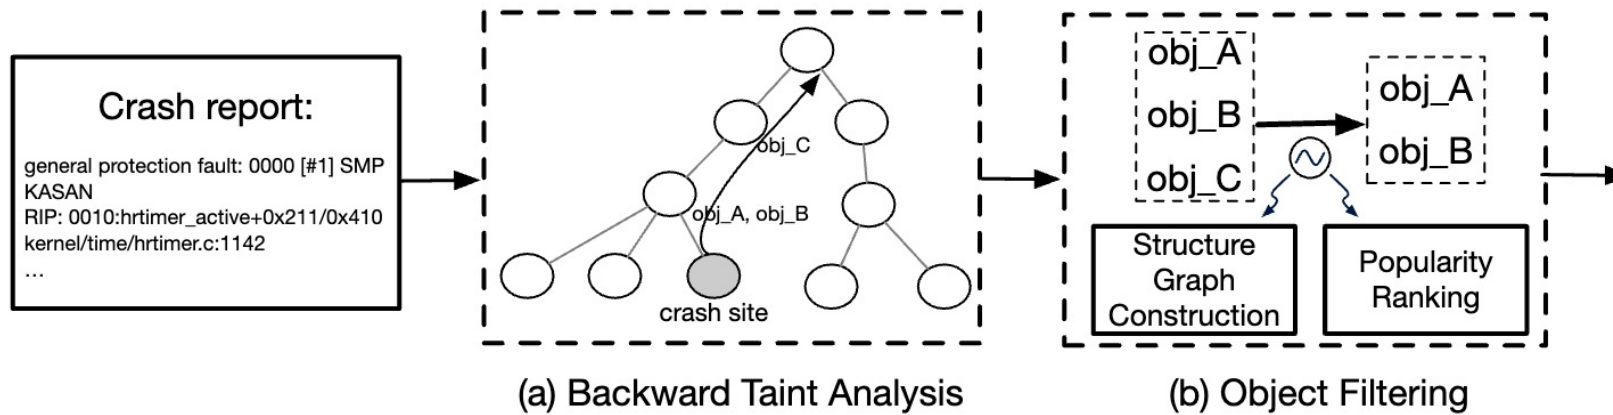


# Why not directed fuzzing

- Identify the root cause of bug is challenge (hard to trigger)
- Only different path is not enough cannot (need vary the context)
- Syzkaller!
- No need to the root casuse
- Syscall sequense to vary the context

# Kernel object fuzzing

- Inappropriate usage of objects
- Incorrect value involved in computation with a kernel object

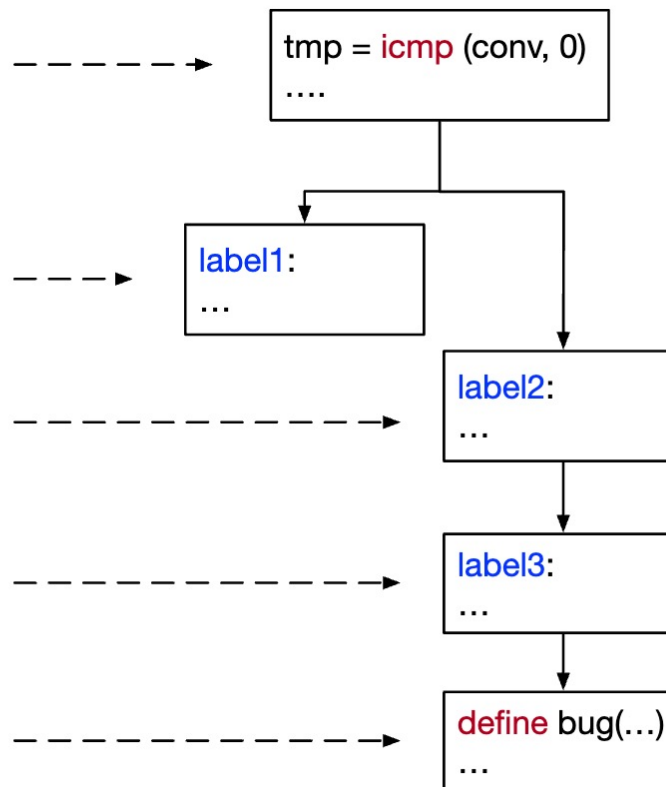


# Critical struct identification

```
1  // in drivers/vhost/vhost.c
2  void vhost_dev_cleanup(struct vhost_dev *dev)
3  {
4      WARN_ON(!list_empty(&dev->work_list));
5      if (dev->worker) {
6          kthread_stop(dev->worker);
7          dev->worker = NULL;
8          dev->kcov_handle = 0;
9      }
10 }
11 // in include/asm-generic/bug.h
12 #define WARN_ON(condition) ({          \
13     int __ret_warn_on = !(condition);    \
14     if (unlikely(__ret_warn_on))          \
15         __WARN();                          \
16     unlikely(__ret_warn_on);              \
17 })
```

# Critical struct identification

```
1 || // comparison
2 || tmp = icmp (conv, 0)
3 || // conditional jump
4 || br (tmp, label1, label2)
5 ||
6 || label1:
7 || call @printk(...) // log
8 ||
9 || label2:
10 || br (label3) // direct jump
11 ||
12 || label3:
13 || call @bug(...) // call
14 ||
15 || define bug(..)
16 || call @printk(...) // log
```



# Critical struct identification

```
1  // source code
2  walk->offset = sg->offset;
3
4  // pseudo binary code after instrumentation
5  kasan_check_read(&sg->offset, sizeof(var));
6  tmp = LOAD(&sg->offset, sizeof(var)); // first access
7  kasan_check_write(&walk->offset, sizeof(var));
8  STORE(tmp, &walk->offset); // second access
```



# Critical struct identification

- Backward taint
- Nested struct and union
- Loop counter

# Critical struct identification

- Tainted value definition
- Syscall entry, interrupt handler
- Entry of function

# Kernel Structure Ranking

- Popular struct such as struct list\_head, struct socket

```
1 // definition of struct sk_buff
2 struct sk_buff {
3     union {
4         struct rb_node rbnode;
5     };
6     ...
7     struct skb_ext *extensions;
8 };
```

# Kernel Structure Ranking

```
1 | static inline void *__skb_push(struct sk_buff *skb, ...)
2 | {
3 |     return skb->data;
4 | }
5 |
6 | int ip6_fraglist_init(...)
7 | {
8 |     struct frag_hdr *fh;
9 |     // type casting from void* to struct frag_hdr*
10 |    fh = __skb_push(skb, sizeof(struct frag_hdr));
11 | }
```

# Object-driven Kernel Fuzzing

- Instrument critical objects related statement
- Corpus: unseen basic block with critical object operation
- Corpus: system call covers more mode and the same system call has critical object operation
- Seed: Not use new system call
- Mutation: new system call in seed corpus

# Object-driven Kernel Fuzzing

- Resource and arguments that system calls operate are necessary for successfully triggering a target kernel bug

```
1 | r0 = openat(...,  
2 | '/dev/dsp1\x00');  
3 | ioctl(r0, ...);  
4 | write(r0, ...);  
5 | read(r0, ...);
```

(a) 7022420

```
1 | // initial PoC: max = -1  
2 | bpf$MAP_CREATE(...,  
3 |     @max=0xffffffffffffffff);  
4 | // exit triggers GFP  
5 | exit(0);
```

(b) 692a8c2

- Group the system call specification templates (resource/usage)
- Only Changed: constant, pointer referencing a memory region, checksum, and resource

# Evaluation

- Bug source: Syzbot -> 50 bugs
  - With PoC program
  - No KMSAN
- Five version kernel (5.6 – 5.10) -> 10 bugs
  - Two recently reproducible bug
- Four VM (7 days)
  - Two GREBE + Syzkaller
  - Two without mutation optimization

# Evaluation

SYZ ID	Critical Structures Identified	Initial Error Behavior	Discovered New Error Behaviors	Time (in hours)			
				T1	T2	T3	T4
bdeea91[23]	aead_instance, crypto_aead, , crypto_spawn, pcrypt_instance_ctx crypto_aead_spawn, crypto_type	WARNING: refcount bug in crypto_mod_get	WARNING: refcount bug in crypto_destroy_tfm	6.69	2.62	0.06	1.25
			KASAN: use-after-free Read in crypto_alg_extsize	-	-	-	83.69
5d3cce3[8]	napi_struct, tun_file	general protection fault in hrtimer_active	KASAN: use-after-free Read in free_netdev	-	-	155.76	30.30
			KASAN: use-after-free Read in netif_napi_add	-	-	77.41	9.08
521a764[24]	ax25_address, nr_sock	WARNING: refcount bug in nr_insert_socket	KASAN: use-after-free Read in release_sock	-	-	0.03	4.39
			KASAN: use-after-free Read in nr_release	-	-	-	20.00
			KASAN: use-after-free Read in nr_insert_socket	-	-	-	0.06
			KASAN: use-after-free Write in nr_insert_socket	-	-	-	126.82
			KASAN: use-after-free Read in lock_sock_nested	-	-	-	18.20
229e0b7[25]	delayed_uprobe	general protection fault in delayed_uprobe_remove	KASAN: use-after-free Read in delayed_uprobe_remove	-	-	3.83	6.66
			KASAN: use-after-free Read in uprobe_mmap	-	-	12.69	4.10
			general protection fault in uprobe_mmap	-	-	-	89.49
			KASAN: use-after-free Read in update_ref_ctr	-	-	-	157.46



<b>SYZ ID</b>	<b>Exploitability Change</b>	<b>SYZ ID</b>	<b>Exploitability Change</b>
d1baeb1 [27]	LL $\rightarrow$ L (2) $\star$	de28cb0 [28]	LL $\rightarrow$ L (5)
8eceaff [29]	LL $\rightarrow$ L (2) $\star$	f56bbe6 [30]	LL $\rightarrow$ L (1)
bb7fa48 [31]	LL $\rightarrow$ L (1)	f0ec9a3 [32]	LL $\rightarrow$ L (1)
d767177 [33]	LL $\rightarrow$ L (2)	5d3cce3 [8]	LL $\rightarrow$ L (2) $\star$
460cc94 [34]	LL $\rightarrow$ L (1)	692a8c2 [12]	LL $\rightarrow$ L (12) $\star$
0df4c1a [35]	LL $\rightarrow$ L (3)	4cf5ee7 [36]	LL $\rightarrow$ L (2)
229e0b7 [25]	LL $\rightarrow$ L (3)	502c872 [37]	LL $\rightarrow$ L (1)
163388d [38]	LL $\rightarrow$ L (1)	b36d7e4 [39]	LL $\rightarrow$ L (1)
bdeea91 [23]	LL $\rightarrow$ L (1)	1fd1d44 [40]	LL $\rightarrow$ L (1)
b9b37a7 [41]	LL $\rightarrow$ L (4)	695527b [42]	LL $\rightarrow$ L (1)
0d93140 [43]	LL $\rightarrow$ L (1)	85fd017 [44]	LL $\rightarrow$ L (4) $\star$
b0e30ab [45]	LL $\rightarrow$ L (1)	6a03985 [46]	LL $\rightarrow$ L (3) $\star$
d5222b3 [47]	LL $\rightarrow$ L (1)	575a090 [48]	LL $\rightarrow$ L (1)
3a6c997 [49]	L $\rightarrow$ L (10)	27ae1ae [50]	L $\rightarrow$ L (1)
cbb2898 [51]	L $\rightarrow$ L (1)	4bf11aa [52]	L $\rightarrow$ L (1)
e4be308 [53]	L $\rightarrow$ L (11)	7022420 [11]	L $\rightarrow$ L (1)
3b7409f [54]	L $\rightarrow$ L (1)	ddaf58b [55]	L $\rightarrow$ L (2)